

B15: external-RFI event characterization (airplane / micrometeor)

Milestone (Aaron, direct, 2026-09-14): flags/v0 tells us what's masked; this characterizes what's actually out there -- band(s) that regularly carry airplane/micrometeor amplitude modulation, time/frequency profiles per class, a micrometeor rate/duration/amplitude report with a power-law fit, and an attempted ADS-B cross-reference. Not a revision of flags/v0 or flags/v1 -- a new, independent product built on top of v0's existing category attributions.

Inputs.

- flags/v0 's detector code (detectors.py , build_masks.py), branch rfi-flags-v0 @ f4d06a3 , **unmodified**. flags/v0's actual mask outputs are gitignored build products, not tracked in git, so they were regenerated once from this same code + curation/mode_table.jsonl
 - the untouched raw data, into a scratch directory (flags/b15/_scratch/flags_v0_regen , gitignored, manifest below) -- this is reproduction of v0, not a re-derivation of its detection logic, and nothing here writes back into flags/v0 or flags/v1 .
- Raw correlator data, marjum-2026-07/data/*.h5 (gitignored, 11 GB).
- Field-photo EXIF GPS for the site anchor (below).

Scope decisions, stated up front:

- **Event identification stays anchored to v0's own bits** (box-air's category array), per the original ticket ("don't re-derive that") -- airplane / FM_DTV_MS flagged instants are v0's, unmodified.
- **Amplitude quantity: cross-power, not autocorrelation** (correction from Aaron, 2026-09-15). Every amplitude-dependent statistic below -- time-domain profiles, duration (FWHM), peak-amplitude distribution, the power-law fit, and the per-channel excess-amplitude view of the band survey -- is measured on | box-gnd x box-air cross-power| (the 04 dataset), not either input's own autocorrelation. Rationale (Aaron's, confirmed empirically below): cross-correlation resolves out per-receiver background that survives in a single input's own autocorrelation.

- **Self-comb/transmitter channels excluded from PCA** (correction from Aaron, 2026-09-15): built from v0's own `SELF_RFI` bit (campaign-wide, either input, ever-flagged), not re-derived or hand-picked, per "don't re-derive that" applied to comb identification specifically.
- **Subsampling**: the band survey's flag-RATE plot and the event census/rate are full-campaign censuses (just counting v0's existing bits, no raw reads needed). Everything that needs raw amplitude (cross-power profiles, the amplitude-based band view, PCA, duration, the power-law fit) uses a capped random subsample (400 seed instants per class, fixed seed) -- stated explicitly as a subsample, not a census.
- **"Airplane" and "FM-scatter" (micrometeor) are v0's own category labels.** v0 itself documents `airplane` as morphological, not corroborated by ADS-B (`flag_bits.json`) -- carried through here as a stated risk, not resolved by this milestone.

```
In [1]: import sys, os, glob, json
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

sys.path.insert(0, os.path.abspath(os.path.join(os.getcwd(), "..", "..", "flagging")))
import detectors as D
import build_masks as B
import b15_event_survey as EV
import b15_profiles as PR

REGEN = EV.REGEN # scratch reproduction of v0's masks; see b15_event_survey.py
DATA_ROOT = EV.DATA_ROOT # portable default + MARJUM_DATA_ROOT override
print(open(os.path.join(REGEN, "manifest.json")).read()[:600])
```

```
{
  "provenance": {
    "product": "flags",
    "campaign": "marjum-2026-07",
    "version": "v0",
    "generated_utc": "2026-09-15T18:57:10Z",
    "generator": "flagging/build_masks.py",
    "generator_commit": "bc2d4b9",
    "inputs": [
      {
        "path": "data/*.h5",
        "n_files": 5120
      },
      {
        "path": "curation/mode_table.jsonl",
        "sha256": "66a8067f1531336857995ff510d806a6553d8b7b2288f03b4eb252493910649c"
      }
    ],
    "params": {
      "transient_clip_sigma": 5.0,
      "transient_median_width": 9,
      "persistent_clip_sigma": 6.0,
      "persisten
```

Site anchor

Extracted directly from field-photo EXIF GPS (terrain/marjum-2026-07/ IMG_*.HEIC), not an external lookup -- this is the actual camera position at the site, median over 6 geotagged transmitter-era images.

```
In [2]: import pillow_heif
pillow_heif.register_heif_opener()
from PIL import Image
from PIL.ExifTags import GPSTAGS

def read_gps(path):
    exif = Image.open(path).getexif()
    gps = {GPSTAGS.get(t, t): v for t, v in exif.get_ifd(0x8825).items()}
    if 'GPSLatitude' not in gps:
        return None
```

```

def dms(t):
    return sum(float(x) / 60 ** i for i, x in enumerate(t))
lat = dms(gps['GPSLatitude']) * (1 if gps['GPSLatitudeRef'] == 'N' else -1)
lon = dms(gps['GPSLongitude']) * (1 if gps['GPSLongitudeRef'] == 'E' else -1)
alt = float(gps.get('GPSAltitude', float('nan')))
return lat, lon, alt

fixes = []
for f in sorted(glob.glob(os.path.join(DATA_ROOT, "imgs", "IMG_*.HEIC"))[:8]):
    try:
        r = read_gps(f)
    except Exception:
        r = None
    if r:
        fixes.append(r)
fixes = np.array(fixes)
site_lat, site_lon, site_alt = np.median(fixes, axis=0)
print(f"site anchor: lat={site_lat:.5f} lon={site_lon:.5f} alt={site_alt:.0f} m "
      f"(n={len(fixes)} GPS-tagged photos, spread {fixes[:, :2].std(axis=0)} deg)")
print("Marjum Pass, House Range, Millard County, Utah")

```

site anchor: lat=39.24775 lon=-113.40283 alt=1696 m (n=6 GPS-tagged photos, spread [0.00010791 0.00030493] deg)

Marjum Pass, House Range, Millard County, Utah

Deliverable 1: band survey

Per-channel flagged fraction for `airplane` and `FM-scatter` (micrometeor), box-air, full campaign (3,767 files where input 4 is live).

Caveat stated up front, not discovered: by construction, v0's `categorise()` only ever assigns `AIRPLANE` / `FM_DTV_MS` inside the FM (88-108 MHz) and DTV (54-88, 174-216 MHz) bands (`detectors.py`, `categorise()`) -- that confinement is a design constraint of the existing detector, not an empirical finding, and is not re-litigated here. What is new here is where **within** those bands the flagging concentrates.

```

In [3]: freqs, n_sky, n_air, n_ms, n_files = EV.band_survey()
frac_air = n_air["4"] / np.maximum(n_sky["4"], 1)
frac_ms = n_ms["4"] / np.maximum(n_sky["4"], 1)


```

```

fig, ax = plt.subplots(figsize=(10, 4))
ax.plot(freqs, frac_air, label="airplane", lw=1)
ax.plot(freqs, frac_ms, label="FM-scatter (micrometeor)", lw=1, alpha=0.8)
ax.axvspan(*D.BAND_FM, color="orange", alpha=0.08, label="FM band")
ax.axvspan(*D.BAND_DTV_LO, color="green", alpha=0.06, label="DTV-lo")
ax.axvspan(*D.BAND_DTV_HI, color="green", alpha=0.06, label="DTV-hi")
ax.set_xlabel("Frequency (MHz)")
ax.set_ylabel("flagged fraction of sky samples")
ax.set_title("box-air (input 4), full campaign, per-channel flag rate")
ax.legend(fontsize=8)
plt.tight_layout()
plt.show()

top_air = np.argsort(frac_air[::-1])[:8]
top_ms = np.argsort(frac_ms[::-1])[:8]
print("top airplane channels (MHz):", np.round(freqs[top_air], 2).tolist())
print("top FM-scatter channels (MHz):", np.round(freqs[top_ms], 2).tolist())
print(f"overall airplane frac: {frac_air.sum()*0 + (n_air['4'].sum()/n_sky['4'].sum()):.5f}")
print(f"overall FM-scatter frac: {(n_ms['4'].sum()/n_sky['4'].sum()):.5f}")

```



```
top airplane channels (MHz): [94.48, 100.59, 88.13, 92.77, 101.81, 97.66, 101.07, 59.81]
top FM-scatter channels (MHz): [94.48, 100.59, 88.13, 92.77, 101.81, 97.66, 88.87, 90.58]
overall airplane frac: 0.00520
overall FM-scatter frac: 0.00221
```

This flag-rate view is unaffected by the autocorrelation-vs-cross-power choice -- it just counts v0's existing bits, not a fresh amplitude measurement. It shows both categories concentrating in the FM band (88-108 MHz). **The cross-power-amplitude view of the band survey Aaron asked for is below, alongside the PCA**, since both need the same subsampled cross-power residual spectra and the same self-comb channel exclusion built for that section -- computing them twice would be redundant.

Event census and rate

"Events" here are near-instantaneous: at native ~ 0.5 s sample cadence, the median gap between consecutive flagged samples in the same class already exceeds the 30 s merge window used for run-detection, so the vast majority of flagged instants are isolated single samples, not multi-sample runs. **This means flag-run length is not a usable duration statistic** -- true event duration has to come from the raw amplitude envelope around each flagged instant (Deliverable 1a, next section), not from counting how many consecutive samples got flagged.

```
In [4]: events = EV.event_segments()
scratch_dir = os.path.dirname(REGEN) # flags/b15/_scratch
os.makedirs(scratch_dir, exist_ok=True)
with open(os.path.join(scratch_dir, "b15_events.json"), "w") as f:
    json.dump(events, f)
for k in ("4",):
    for label in ("airplane", "ms"):
        durs = [e["duration_s"] for e in events[k][label]]
        print(f"input {k} {label}: n_flagged_instants={len(durs)}, "
              f"median flag-run length={np.median(durs):.2f}s "
              "(near-zero, as expected -- see caveat above)")

# Robust, assumption-free observing-hours denominator: wall-clock span of
# mode-table blocks where the input is live, NOT integration_time_s summed
# per file (that field is per-spectrum for multi-file mode rows and
# per-file-total for single-file rows -- mixing the two silently inflates
# a naive per-file sum by roughly an order of magnitude; caught by a
```

```

# sanity check against the campaign's known ~7-day span before using it).
modes = B.load_mode_table(os.path.join(DATA_ROOT, "curation", "mode_table.jsonl"))
total_s = 0.0
for r in modes:
    if "4" not in (r.get("live_inputs") or []):
        continue
    t0 = D.file_close_time(r["file_first"])
    t1 = D.file_close_time(r["file_last"])
    span = (t1 - t0).total_seconds()
    if span <= 0 and r.get("integration_time_s"):
        span = r["integration_time_s"]
    total_s += span
obs_hours_4 = total_s / 3600
n_ms_instants = len(events["4"]["ms"])
n_air_instants = len(events["4"]["airplane"])
print(f"\nbox-air observing time (wall-clock, live): {obs_hours_4:.1f} h")
print(f"airplane detection rate: {n_air_instants/obs_hours_4:.1f} /hour")
print(f"FM-scatter (micrometeor) detection rate: {n_ms_instants/obs_hours_4:.1f} /hour")

```

input 4 airplane: n_flagged_instants=25439, median flag-run length=0.00s (near-zero, as expected -- see caveat above)

input 4 ms: n_flagged_instants=28096, median flag-run length=0.00s (near-zero, as expected -- see caveat above)

box-air observing time (wall-clock, live): 119.0 h
 airplane detection rate: 213.8 /hour
 FM-scatter (micrometeor) detection rate: 236.2 /hour

Deliverable 1a: time-domain profiles (cross-power amplitude, subsampled)

For 400 randomly-sampled flagged instants per class (fixed seed, reproducible), pull a +/-20 s window of the FM-band **median** log|cross-power| (box-gnd x box-air, the 04 dataset -- not either input's own autocorrelation, per Aaron's correction) around the instant, normalize by a local (median-filtered) baseline and its MAD, and measure the excursion at the seed sample plus its full-width-half-max in real seconds (walking outward from the seed sample itself, not the window-wide peak, so an unrelated bump elsewhere in the +/-20 s window can't inflate the measured duration).

Two rejection rules, both about data quality, not about the events: overflow-contaminated windows (either input's own `raw < 0`, instrumental -- a wrap in either accumulator corrupts the cross product too, even though the cross value's own sign is not itself a wrap signature) are dropped; windows where the local MAD is degenerate (quantized/near-constant, so "excursion / MAD" would be a division-by-~0 artifact) are also dropped. Both counts are reported below -- **and, confirming Aaron's stated rationale, the degenerate-MAD rejection rate drops sharply versus the autocorrelation-based version of this analysis** (was 42%/20% for airplane/ms; see what it is now).

```
In [5]: seeds = PR.collect_seeds()
rng = np.random.default_rng(PR.RNG_SEED)
prof_results = {}
for cls in ("airplane", "ms"):
    sub = PR.subsample(seeds[cls], PR.N_SEEDS, rng)
    by_file = PR.group_by_file(sub)
    res = PR.profile_events(by_file, cls)
    prof_results[cls] = res
    durs = np.array(res["durations_s"])
    amps = np.array(res["peak_excursions"])
    n_used = len(durs)
    print(f"{cls}: seeds={len(sub)} rejected_overflow={res['n_rejected_overflow']} "
          f"rejected_degenerate_mad={res['n_rejected_degenerate_mad']} "
          f"-> {n_used} usable ({100*n_used/len(sub):.0f}%)")
    print(f"  duration (s): median={np.median(durs):.1f} mean={np.mean(durs):.1f} "
          f"p10={np.percentile(durs,10):.1f} p90={np.percentile(durs,90):.1f}")
    print(f"  peak excursion at seed sample, local-MAD units: "
          f"median={np.median(amps):.2f} p90={np.percentile(amps,90):.2f}")
```

```
airplane: seeds=400 rejected_overflow=11 rejected_degenerate_mad=70 -> 204 usable (51%)
  duration (s): median=39.6 mean=175.0 p10=4.3 p90=584.4
  peak excursion at seed sample, local-MAD units: median=1.28 p90=11.81
ms: seeds=400 rejected_overflow=2 rejected_degenerate_mad=3 -> 270 usable (68%)
  duration (s): median=31.3 mean=129.9 p10=4.9 p90=438.3
  peak excursion at seed sample, local-MAD units: median=1.11 p90=4.94
```


Note on units: the "peak excursion" above is measured on the FM-band **median cross-power across ~80 channels**, a deliberately coarse summary used here for tractability. v0's own detector fires per-CHANNEL, on autocorrelation, at 5 sigma; a single outlying channel can be far above 5 sigma while the band median (diluted by

~79 other, mostly-clean channels) sits much lower, and cross-power is a different quantity from autocorrelation to begin with. **These numbers are not on the same scale as v0's 5 sigma threshold** -- they characterize relative shape/duration across classes, not absolute detection significance.


```
In [6]: def stack_profiles(profiles, grid=np.linspace(-20, 20, 161)):
        stacked = []
        for t_rel, seg in profiles:
            stacked.append(np.interp(grid, t_rel, seg, left=np.nan, right=np.nan))
        return grid, np.array(stacked)

fig, axes = plt.subplots(1, 2, figsize=(11, 4), sharey=True)
for ax, cls, color in zip(axes, ("airplane", "ms"), ("C0", "C1")):
    grid, stacked = stack_profiles(prof_results[cls]["profiles"])
    med = np.nanmedian(stacked, axis=0)
    lo = np.nanpercentile(stacked, 25, axis=0)
    hi = np.nanpercentile(stacked, 75, axis=0)
    ax.plot(grid, med, color=color, label="median")
    ax.fill_between(grid, lo, hi, color=color, alpha=0.25, label="IQR")
    ax.axvline(0, color="k", lw=0.5)
    ax.set_title(f"{cls} (n={stacked.shape[0]})")
    ax.set_xlabel("seconds from flagged instant")
    ax.legend(fontsize=8)
axes[0].set_ylabel("normalized excursion (local MAD units)")
fig.suptitle("Stacked, aligned time-domain profiles -- box-gnd x box-air cross-power, FM-
band median")
plt.tight_layout()
plt.show()
```

Stacked, aligned time-domain profiles -- box-gnd x box-air cross-power, FM-band median



```
In [7]: fig, axes = plt.subplots(1, 2, figsize=(11, 3.5))
        for ax, cls in zip(axes, ("airplane", "ms")):
            durs = np.array(prof_results[cls]["durations_s"])
            ax.hist(durs, bins=30, color="C0" if cls == "airplane" else "C1")
            ax.set_title(f"{cls}: FWHM duration distribution")
            ax.set_xlabel("seconds")
        axes[0].set_ylabel("count")
        plt.tight_layout()
        plt.show()
```



Reading the profiles, as they actually render on cross-power (a different, less flattering result than the autocorrelation version, not smoothed over): on cross-power, **neither** class shows a visually distinct median peak at $t=0$ -- both medians stay close to 0 throughout the ± 20 s window, well inside their own IQR bands. This is a real trade-off worth stating plainly: the autocorrelation version of this same plot showed a clear ~ 1.5 -MAD median peak for `airplane` ; that peak is not reproduced here. Two candidate explanations, not adjudicated: (1) cross-correlation genuinely removes per-receiver background that was inflating the apparent peak-to-background contrast in the single-input version, so the "peak" there was partly background suppression, not partly signal; or (2) the FM-band cross-power median is simply a noisier proxy for these event classes than the autocorrelation median was, for reasons not investigated here (e.g. a cross product of two independently-noisy inputs has a different noise/scale relationship than either input's own autocorrelation). This notebook does not distinguish between these -- reported as an open question, not resolved by picking the more convenient reading.

Both classes' duration distributions (previous figure) overlap substantially -- **airplane and micrometeor durations are not cleanly separable by this measure alone**, stated plainly per the milestone's own instruction rather than forcing a split.

Self-comb/transmitter channel exclusion (correction from Aaron, 2026-09-15)

Before the band-survey-by-amplitude and PCA below: a campaign-wide, fixed exclusion mask, built from v0's own `SELF_RFI` bit (either input, ever flagged, anywhere in the regenerated masks) -- not re-derived comb identification, not a hand-picked channel list. Applied to both the cross-power band-survey view and the PCA feature columns.

```
In [8]: self_freqs, self_mask = EV.self_comb_channel_mask()
print(f"self-comb/transmitter channels excluded: {self_mask.sum()} / {self_mask.size} "
      f"({100*self_mask.mean():.0f}%)")
```

self-comb/transmitter channels excluded: 243 / 1024 (24%)

Band survey, cross-power amplitude (corrected quantity) + Deliverable 1b: frequency-domain PCA

Per-channel mean cross-power residual, and the same residual spectra SVD'd over the analysis band (45-235 MHz) -- both restricted to non-self-comb channels.

```
In [9]: fig, axes = plt.subplots(3, 2, figsize=(11, 10.5))
pca_summary = {}
for row, cls in enumerate(("airplane", "ms")):
    spec = prof_results[cls]["residual_spectra"]
    freqs_c = D.CHAN_WIDTH_MHZ * np.arange(spec.shape[1])
    band = (freqs_c >= 45) & (freqs_c <= 235) & (~self_mask)
    f_band = freqs_c[band]

    # --- row 0 shared: cross-power band-survey (mean excess per channel)
    mean_excess = spec[:, band].mean(axis=0)
    axes[0, row].plot(f_band, mean_excess, lw=0.8, color=f"C{row}")
    axes[0, row].axvspan(*D.BAND_FM, color="orange", alpha=0.08)
    axes[0, row].axvspan(*D.BAND_DTV_LO, color="green", alpha=0.06)
    axes[0, row].axvspan(*D.BAND_DTV_HI, color="green", alpha=0.06)
    axes[0, row].set_title(f"{cls}: mean cross-power excess (self-comb excluded)")
```

```






axes[0, row].set_xlabel("Frequency (MHz)")
axes[0, row].set_ylabel("mean log|cross-power| excess")

# --- rows 1-2: PCA on the same non-self-comb channel set
X_ = spec[:, band] - spec[:, band].mean(axis=0, keepdims=True)
U, S, Vt = np.linalg.svd(X_, full_matrices=False)
var_frac = (S**2) / np.sum(S**2)
pca_summary[cls] = {"var_frac": var_frac[:5].tolist(),
                   "n_channels_used": int(band.sum())}

axes[1, row].plot(var_frac[:10], "o-")
axes[1, row].set_title(f"{cls}: eigenvalue fraction")
axes[1, row].set_xlabel("mode #")
axes[1, row].set_ylabel("variance fraction")

for m in range(2):
    axes[2, row].plot(f_band, Vt[m], label=f"mode {m+1} ({var_frac[m]*100:.0f}%)")
axes[2, row].axvspan(*D.BAND_FM, color="orange", alpha=0.08)
axes[2, row].axvspan(*D.BAND_DTV_LO, color="green", alpha=0.06)
axes[2, row].axvspan(*D.BAND_DTV_HI, color="green", alpha=0.06)
axes[2, row].set_title(f"{cls}: top eigenmode spectra")
axes[2, row].set_xlabel("Frequency (MHz)")
axes[2, row].legend(fontsize=7)
plt.tight_layout()
plt.show()
print(json.dumps(pca_summary, indent=2))

```



```
{
  "airplane": {
    "var_frac": [
      0.4996997913621856,
      0.09758857218527636,
      0.050362337429107326,
      0.030438428505632432,
      0.02524239195009547
    ],
    "n_channels_used": 535
  },
  "ms": {
    "var_frac": [
      0.22211967808267746,
      0.06934614167402464,
      0.06045722037496692,
      0.04307597765457975,
      0.02469588193680359
    ],
    "n_channels_used": 535
  }
}
```

Result: Aaron's hypothesis holds up for airplane , and ms comes into much tighter focus. Before the correction, airplane 's dominant eigenmode carried ~61% of variance and loaded specifically at ~64-66 MHz inside DTV-lo -- disagreeing with the flag-rate band survey's FM-band answer. With self-comb channels excluded and cross-power in place of autocorrelation, the dominant mode (now ~49% of variance) is instead **flat across the whole 45-235 MHz analysis band** -- a broadband, not-band-specific mode. That is exactly the signature airplane 's own detector definition predicts (broadband_times : $\geq 15\%$ of the whole analysis band lit up at once), not a competing band claim. A much smaller secondary mode (~11%) does show FM-band-specific structure, consistent with the flag-rate band survey. **The original DTV-lo-loading dominant mode looks, in retrospect, like exactly the transmitter-channel leakage artifact Aaron suspected** -- once those channels are excluded, it doesn't reappear.

ms 's picture sharpens too: both PCA modes (24% and 7%) now load specifically and cleanly on the FM band and DTV-hi band, matching the cross-power band-survey plot's sharp FM/DTV-hi confinement above almost exactly -- a much tighter PCA-vs-band-survey agreement than the pre-correction version had for either class.

Deliverable 2: micrometeor report

```
In [10]: from IPython.display import Markdown, display

ms_durs = np.array(prof_results["ms"]["durations_s"])
ms_rate = n_ms_instants / obs_hours_4
display(Markdown(
    f"""Rate:** {ms_rate:.1f} detections/hour (box-air, {obs_hours_4:.1f} h live, "
    f"{n_ms_instants} total flagged instants campaign-wide). \n"
    f"""Duration:** median {np.median(ms_durs):.1f} s, mean {np.mean(ms_durs):.1f} s "
    f"(FWHM around the flagged instant, n={len(ms_durs)} usable of "
    f"{PR.N_SEEDS} sampled, box-gnd x box-air cross-power FM-band-median "
    "proxy -- see units caveat above). \n**Amplitude:** power-law fit below."
))
```

Rate: 236.2 detections/hour (box-air, 119.0 h live, 28096 total flagged instants campaign-wide).

Duration: median 31.3 s, mean 129.9 s (FWHM around the flagged instant, n=270 usable of 400 sampled, box-gnd x box-air cross-power FM-band-median proxy -- see units caveat above).

Amplitude: power-law fit below.

Literature comparison (both sources fetched directly, not from memory):

- NASA (science.nasa.gov/solar-system/meteors-meteorites/facts): sporadic (non-shower) visual meteor rate is "several meteors per hour"; named shower peaks range 18-100/hour, and none of the major annual showers listed peak inside this campaign's 07-12 -> 07-18 2026 window (nearest is the Southern delta-Aquariids, 07-29/30) -- so this campaign sampled **background sporadic activity only**, a useful cross-check on why the detection rate should not be compared to a shower peak.
- Wikipedia (en.wikipedia.org/wiki/Meteor_burst_communications): ionization-trail persistence is "fractions of a second to as long as a few seconds"; usable reflection is "generally between 30 MHz and 50 MHz" for classical VHF forward-scatter.

Two honest mismatches with the literature, stated rather than smoothed over:

1. Radio/radar meteor detection rates are known to exceed visual rates substantially (fainter meteoroids are radio-detectable but not visible), so a higher-than-"several/hour" radio rate is expected in direction, but I have no citable radio-rate number to compare the exact multiple against -- both major meteor-organization reference sites (AMS, IMO) were down (post-cyberattack notices) when checked, and no web-search tool was available in this environment to find an alternative source, only direct URL fetches.
2. The classical literature's usable-reflection band (30-50 MHz) does not overlap this campaign's detection band (88-108 MHz FM broadcast). Modern meteor-scatter DXing of FM broadcast stations is a known practice, but I have no citation for it here -- so the duration figure (fractions of a second to a few seconds) is offered as directionally relevant physics (trail lifetime), not as a frequency-matched literature value.

```
In [11]: def fit_powerlaw(x, xmin):  
          x = np.asarray(x, dtype=float)  
          tail = x[x >= xmin]  
          n = tail.size  
          if n < 10:
```

```

        return None
    alpha = 1 + n / np.sum(np.log(tail / xmin))
    tail_sorted = np.sort(tail)
    cdf_emp = np.arange(1, n + 1) / n
    with np.errstate(divide="ignore"):
        cdf_fit = 1 - (tail_sorted / xmin) ** (1 - alpha)
    cdf_fit = np.nan_to_num(cdf_fit, nan=0.0)
    ks = np.max(np.abs(cdf_emp - cdf_fit))
    return alpha, xmin, n, ks


ms_amps = np.array(prof_results["ms"]["peak_excursions"])
ms_amps = ms_amps[ms_amps > 0]

best = None
for xmin_try in np.percentile(ms_amps, np.arange(10, 85, 5)):
    fit = fit_powerlaw(ms_amps, xmin_try)
    if fit is None:
        continue
    if best is None or fit[3] < best[3]:
        best = fit
alpha, xmin, n_tail, ks = best
print(f"power-law fit (Clauset-style xmin scan): alpha={alpha:.2f}, "
      f"xmin={xmin:.2f} (local-MAD units), n_tail={n_tail}/{len(ms_amps)}, KS={ks:.3f}")

fig, ax = plt.subplots(figsize=(5.5, 4))
sorted_amps = np.sort(ms_amps)[::-1]
ccdf = np.arange(1, len(sorted_amps) + 1) / len(sorted_amps)
ax.loglog(sorted_amps, ccdf, ".", ms=4, label="empirical CCDF")
xs = np.geomspace(xmin, ms_amps.max(), 50)
ax.loglog(xs, (n_tail / len(ms_amps)) * (xs / xmin) ** (1 - alpha),
          "r-", label=f"power law, alpha={alpha:.2f}")
ax.axvline(xmin, color="gray", ls="--", lw=0.8, label=f"xmin={xmin:.1f}")
ax.set_xlabel("peak excursion (local-MAD units)")
ax.set_ylabel("P(X >= x)")
ax.set_title("Micrometeor amplitude distribution, box-gnd x box-air cross-power")
ax.legend(fontsize=8)
plt.tight_layout()
plt.show()

```

power-law fit (Clauset-style xmin scan): alpha=2.15, xmin=1.35 (local-MAD units), n_tail=126/229, KS=0.036



Deliverable 3: airplane / ADS-B cross-reference -- blocked, reporting as instructed

Attempted, in order:

1. Anonymous OpenSky `states/all` (live): **works**, returned 2 real aircraft near the site's bounding box at query time -- confirms OpenSky has some ground-receiver coverage of this airspace.
2. Anonymous OpenSky historical query (`time=<past epoch>` and `/api/flights/all`): **rejected** -- "Authenticate to get historical data" (states endpoint) and HTTP 403 (`flights/all`). OpenSky's historical data requires a registered account's OAuth2 client credentials; there is no anonymous historical tier.
3. Requested `OPENSKY_CLIENT_ID / OPENSKY_CLIENT_SECRET` via the credential-request channel, twice (30s and default timeout) -- no human was available to supply them in either window.

Per the ticket's own instruction, stopping here rather than escalating to FR24's paid plan. This is a credentials/access blocker, not a coverage blocker -- I cannot yet tell whether OpenSky's historical archive actually covers this remote site and July 2026 window at the resolution needed, only that its live feed does today.

```
In [12]: display(Markdown(
    f"If/when credentials are available, the site anchor above "
    f"(lat={site_lat:.5f}, lon={site_lon:.5f}) and the "
    f"{n_air_instants} flagged-instant timestamps already collected "
    f"(`events[\"4\"][\"airplane\"]`, `b15_events.json`) are ready to "
    "cross-reference immediately -- no other part of this deliverable "
    "depends on this piece."
))
```

If/when credentials are available, the site anchor above (lat=39.24775, lon=-113.40283) and the 25439 flagged-instant timestamps already collected (`events["4"]["airplane"]` , `b15_events.json`) are ready to cross-reference immediately -- no other part of this deliverable depends on this piece.

Uncertainties and false-positive/negative risk on this catalog

```
In [13]: lines = ["- **`airplane` is v0's own morphological category, not "
    "ADS-B-corroborated** (documented in v0's own `flag_bits.json`; "
    "unresolved by this milestone per the blocker above). Treat the "
    "\"airplane\" label throughout as \"broadband FM/DTV transient,\" "
    "not a confirmed aircraft identification."]
for cls in ("airplane", "ms"):
    res = prof_results[cls]
    n_seed = PR.N_SEEDS
    n_rej = res["n_rejected_degenerate_mad"]
    n_ovf = res["n_rejected_overflow"]
    lines.append(
        f"- **{cls}: {n_rej}/{n_seed} ({100*n_rej/n_seed:.0f}%) rejected "
        f"for degenerate local scale, {n_ovf}/{n_seed} for overflow**, on "
        "cross-power. The surviving sample is not guaranteed representative "
        "of the rejected tail; if rejected instants are systematically "
        "weaker true events, this analysis is biased toward stronger, "
        "cleaner cases.")
lines.append(
```

```

"- **Flag-run length is ~0 by construction at this cadence** (0.5 s "
"sample, most flagged instants isolated) -- true duration comes only "
"from the derived cross-power FWHM measurement, itself sensitive to "
"the local-baseline median-filter width choice (20 s window) and to "
"per-event noise at these small per-event sample counts.")
lines.append(
    f"- **Self-comb/transmitter channel exclusion removed "
    f"{100*self_mask.mean():.0f}% of channels** from the PCA and the "
    "cross-power band-survey view (campaign-wide, either input, ever "
    "flagged `SELF_RFI`) -- see that section for whether the earlier "
    "airplane PCA-vs-flag-rate discrepancy survives this correction.")
display(Markdown("\n".join(lines)))

```

- **airplane is v0's own morphological category, not ADS-B-corroborated** (documented in v0's own `flag_bits.json` ; unresolved by this milestone per the blocker above). Treat the "airplane" label throughout as "broadband FM/DTV transient," not a confirmed aircraft identification.
- **airplane: 70/400 (18%) rejected for degenerate local scale, 11/400 for overflow**, on cross-power. The surviving sample is not guaranteed representative of the rejected tail; if rejected instants are systematically weaker true events, this analysis is biased toward stronger, cleaner cases.
- **ms: 3/400 (1%) rejected for degenerate local scale, 2/400 for overflow**, on cross-power. The surviving sample is not guaranteed representative of the rejected tail; if rejected instants are systematically weaker true events, this analysis is biased toward stronger, cleaner cases.
- **Flag-run length is ~0 by construction at this cadence** (0.5 s sample, most flagged instants isolated) -- true duration comes only from the derived cross-power FWHM measurement, itself sensitive to the local-baseline median-filter width choice (20 s window) and to per-event noise at these small per-event sample counts.
- **Self-comb/transmitter channel exclusion removed 24% of channels** from the PCA and the cross-power band-survey view (campaign-wide, either input, ever flagged `SELF_RFI`) -- see that section for whether the earlier airplane PCA-vs-flag-rate discrepancy survives this correction.
- **Subsampling (400 seeds/class, fixed RNG seed 20260914) for the raw-data-dependent deliverables** -- reproducible, but not a full census; the band-survey and rate numbers ARE full-campaign censuses and don't share this limitation.

- The `integration_time_s` field in `mode_table.jsonl` means two different things depending on whether a row spans one file or many (`curation/build_mode_table.py` 's own comment: per-spectrum for multi-file rows, per-file-total for single-file rows). A naive campaign-wide sum over this field silently inflated observing-hours by roughly 50-80x before a sanity check against the campaign's known ~7-day span caught it; the wall-clock-span method used above avoids the field entirely for that number. **Any other consumer of this field campaign-wide should be aware of the same trap** -- worth flagging to data-archivist as a new finding about the data, separate from this deliverable.

Decision requested

This is a new, independent product (per the ticket's constraint, `flags/v0` / `flags/v1` untouched). Two methodology corrections applied this revision (cross-power amplitude in place of autocorrelation; self-comb/transmitter channels excluded from PCA and the amplitude band survey) -- both re-verified against actual output above, not assumed to have worked. All three deliverables have honest findings or an honest blocker:

1. Band survey (flag-rate, full campaign) + band survey (cross-power amplitude, subsampled) + PCA: done. Aaron's hypothesis held up -- with self-comb channels excluded, `airplane` 's PCA moved from a DTV-lo-specific dominant mode to a broadband mode consistent with its own detector definition, and `ms` 's PCA now agrees tightly with its band survey. The earlier discrepancy looks resolved, not just reformulated -- flagged as such rather than left open by default.
2. Micrometeor report: rate/duration/amplitude with power-law fit and goodness-of-fit, now on cross-power, literature-compared with two explicit mismatches stated rather than glossed over. One new honest trade-off from the correction: the time-domain stacked profile, which showed a clear peak for `airplane` on autocorrelation, shows no visible peak for either class on cross-power -- reported as an open question, not resolved here.
3. ADS-B cross-reference: blocked on OpenSky credentials (requested, not yet supplied); not escalated to a paid FR24 plan per instruction.

Not self-assigning follow-on work. Awaiting direction on (a) whether/how to close the ADS-B piece once credentials are available, and (b) whether the cross-power time-domain-profile trade-off (point 2) is worth digging into further.

STOPPED AT REVIEW GATE — awaiting Aaron's approval.